

Constrained optimization

A typical optimization problem under constraints, look like:

$$\min_{\mathbf{x}} f(\mathbf{x}) \text{ such that } g_1(\mathbf{x}) \leq 0, g_2(\mathbf{x}) \leq 0, g_3(\mathbf{x}) \geq 0, \text{ etc.},$$

Lagrange multiplier methods are one of the most popular methods to solve such minimization problem. The objective function $J = f(\mathbf{x})$ is Lagrangian as by the constraint equations through a set of non-negative multiplicative Lagrange multipliers, $\lambda_j \geq 0$. The Lagrangian as objective function, $J_A(\mathbf{x}, \lambda_1, \lambda_2, \lambda_3)$, is,

$$f(\mathbf{x}) + \lambda_1 g_1(\mathbf{x}) + \lambda_2 g_2(\mathbf{x}) - \lambda_3 g_3(\mathbf{x}).$$

Notice, that the last term is negative is due to the \geq inequality. The reason is similar to the motivation around the penalty-induced regression methods. Specifically, since the aim is to minimize $f(\mathbf{x})$, the idea is to keep the Lagrangian as objective function a minimization problem. Since $g_1(\mathbf{x}) \leq 0, g_2(\mathbf{x}) \leq 0, g_3(\mathbf{x}) \geq 0$ for all \mathbf{x} , we have $J_A(\mathbf{x}, \lambda_1, \lambda_2, \lambda_3) \leq f(\mathbf{x})$.

Thus, $\max_{\lambda_1, \lambda_2, \lambda_3} J_A(\mathbf{x}, \lambda_1, \lambda_2, \lambda_3) \leq f(\mathbf{x})$ as $J_A(\mathbf{x}, \lambda_1, \lambda_2, \lambda_3) \leq f(\mathbf{x})$ holds for all positive λ_i 's. And our optimization problem becomes,

$$\max_{\lambda_1, \lambda_2, \lambda_3} \min_{\mathbf{x}} J_A(\mathbf{x}, \lambda_1, \lambda_2, \lambda_3) \text{ such that } \lambda_1, \lambda_2, \lambda_3 \geq 0.$$

1. For minimization problem, we need to make sure $J_A(\mathbf{x}, \boldsymbol{\lambda}) \leq f(\mathbf{x})$ for $\lambda \geq 0$.
2. For maximization problem, we need to make sure $J_A(\mathbf{x}, \boldsymbol{\lambda}) \geq f(\mathbf{x})$ for $\lambda \geq 0$.

A popular strategy to solve Lagrange multiplier based constraint optimization is,

1. First, compute the derivate with respect to \mathbf{x} as $\frac{\partial J_A(\mathbf{x}, \lambda_1, \lambda_2, \lambda_3)}{\partial \mathbf{x}}$
2. Also the derivatives with respect to λ_j 's which are $g_1(\mathbf{x})$, $g_2(\mathbf{x})$ and $g_3(\mathbf{x})$.

These equations define the relations to be used to derive expressions for, or compute values of, the value \mathbf{x}^* (the optimal value). However, the optimal value of some λ_i 's may be zero (thus that constraint will not have any impact on the optimization) if the solution leads to something inconsistent.

We first (blindly) solve the system of equations that we get after taking first derivatives with respect of \mathbf{x} and λ_j 's. In case of multiple constraints, we get multiple sets of possible solutions. If an inequality $g_j(\mathbf{x}) \geq 0$ constrains the optimum point, the corresponding Lagrange multiplier, λ_j^* , is positive. If an inequality $g_j(\mathbf{x}) \geq 0$ does not constrain the optimum point, the corresponding Lagrange multiplier, λ_j^* , is set to zero.

Consider the following constrained minimization problem:

$$\min_x \frac{1}{2}kx^2, \text{ for } k > 0 \text{ such that } x \geq b$$

The inequality is like $g_3()$ in our earlier definition. Thus, The Lagrangian as objective function is,

$$\max_{\lambda} \min_x \frac{1}{2}kx^2 - \lambda(x - b)$$

Now the taking derivatives would give us two equations $kx - \lambda = 0$ and $x - b = 0$. Let (x^*, λ^*) are the optimal values.

Case 1: $b > 0$ Then $x = b$ and $\lambda = kb > 0$ which satisfies the required condition. This is thus an optimal solution for x and λ .

Case 2: $b < 0$ Then, $x = b$ and $\lambda = kb < 0$, does not satisfy the required positivity condition on λ . Hence, we set $\lambda^* = 0$. Having set $\lambda^* = 0$, we have to minimize $\frac{1}{2}kx^2$ which leads to $x^* = 0$.

Hence, when $b > 0$, we have $(x^*, \lambda^*) = (b, kb)$, but when $b < 0$, we have $(x^*, \lambda^*) = (0, 0)$.

In this example, we can verify following two important properties:

1. $J_A(x^*, \lambda) \leq J_A(x^*, \lambda^*) \leq J_A(x, \lambda^*)$. Thus the Lagrangian as objective function is minimized at $x = x^*$ when we set $\lambda = \lambda^*$ and is maximized at $\lambda = \lambda^*$ when we set $x = x^*$.
2. Formally, $\min_x J_A(x, \lambda^*) = \max_{\lambda} J_A(x^*, \lambda) = J_A(x^*, \lambda^*)$.

Short digression

Minimax inequality It can be shown easily that $\max_{\lambda} \min_x J_A(x, \lambda) \leq \min_x \max_{\lambda} J_A(x, \lambda)$.

Proof. $J_A(x, \lambda) \leq \max_{\lambda} J_A(x, \lambda)$ for any x and λ . Also, $J_A(x, \lambda) \geq \min_x J_A(x, \lambda)$ for any x and λ . Combining the two, we have $\min_x J_A(x, \lambda) \leq \max_{\lambda} J_A(x, \lambda)$. Now, take the minimum over x on the right-hand side, then the maximum over λ on the left-hand side. \square

Let $u^* = \min_x \max_{\lambda} J_A(x, \lambda)$ and $l^* = \max_{\lambda} \min_x J_A(x, \lambda)$. The difference $u^* - l^*$ is called the duality gap.

Minimax equality It can be shown that the gap $u^* - l^*$ is zero if:

- The domains of x and λ are both convex, and one of them is compact.
- The function J_A is convex-concave: $J_A(\cdot, \lambda)$ is convex for any given λ , and $J_A(x, \cdot)$ is concave for any given x .
- The function J_A is continuous

Above strategy to solve Lagrangian as objective function may become hard when we are in a multiparameter setting with several constraints. We then take another way which is the following,

1. First, compute $\frac{\partial J_A(\mathbf{x}, \lambda_1, \lambda_2, \lambda_3)}{\partial \mathbf{x}} = 0$ and solve \mathbf{x} in terms of λ . Let the solution be, $h(\lambda_1, \lambda_2, \lambda_3)$
2. Use this solution of x and solve $\max_{\lambda_1, \lambda_2, \lambda_3} J_A(h(\lambda_1, \lambda_2, \lambda_3), \lambda_1, \lambda_2, \lambda_3)$

This strategy leads to a new objective function $J'_A(\lambda_1, \lambda_2, \lambda_3) = J_A(h(\lambda_1, \lambda_2, \lambda_3), \lambda_1, \lambda_2, \lambda_3)$ for our original problem. This is also called the ‘dual’ problem. Whenever we write the optimization problem in terms of the Lagrange multiplier’s only, it is called the *Dual*. Solving a dual problem is easier due to fewer number of parameters in the objective function. If I am too strict, this is called Lagrange dual. There are other types of dual representations as well.

Example:

$$\min_{\mathbf{x}} \frac{1}{2} \mathbf{x}^T \mathbf{M} \mathbf{x} + \mathbf{c}^T \mathbf{x} + d \text{ such that } \mathbf{A} \mathbf{x} \leq \mathbf{b}.$$

First, we write the Lagrangian as objective function, which is (*the primal form*),

$$\max_{\boldsymbol{\lambda}} \min_{\mathbf{x}} J(\mathbf{x}, \boldsymbol{\lambda}) = \frac{1}{2} \mathbf{x}^T \mathbf{M} \mathbf{x} + \mathbf{c}^T \mathbf{x} + d + \boldsymbol{\lambda}^T (\mathbf{A} \mathbf{x} - \mathbf{b}),$$

such that $\boldsymbol{\lambda} \geq 0$.

We first solve \mathbf{x} in terms of $\boldsymbol{\lambda}$ from $\frac{\partial J(\mathbf{x}, \boldsymbol{\lambda})}{\partial \mathbf{x}} = 0 \implies \mathbf{M} \mathbf{x} + \mathbf{c} + \mathbf{A}^T \boldsymbol{\lambda} = 0 \implies \mathbf{x} = -\mathbf{M}^{-1}(\mathbf{c} + \mathbf{A}^T \boldsymbol{\lambda})$.

Hence, the Lagrangian dual problem (*the dual form*) is,

$$\max_{\boldsymbol{\lambda}} \frac{1}{2} (\mathbf{c} + \mathbf{A}^T \boldsymbol{\lambda})^T \mathbf{M}^{-1} (\mathbf{c} + \mathbf{A}^T \boldsymbol{\lambda}) - \mathbf{c}^T \mathbf{M}^{-1} (\mathbf{c} + \mathbf{A}^T \boldsymbol{\lambda}) + d + \boldsymbol{\lambda}^T (\mathbf{A} \{-\mathbf{M}^{-1}(\mathbf{c} + \mathbf{A}^T \boldsymbol{\lambda})\} - \mathbf{b}).$$

Upon simplification,

$$\max_{\boldsymbol{\lambda}: \boldsymbol{\lambda} \geq 0} -\frac{1}{2} \boldsymbol{\lambda}^T \mathbf{A} \mathbf{M}^{-1} \mathbf{A}^T \boldsymbol{\lambda} - (\mathbf{c}^T \mathbf{M}^{-1} \mathbf{A}^T + \mathbf{b}^T) \boldsymbol{\lambda} - \frac{1}{2} \mathbf{c}^T \mathbf{M}^{-1} \mathbf{c} + d.$$

Note that if the quadratic term in the primal form (quadratic in the primal variable \mathbf{x}) is positive, then the quadratic term in the dual form (quadratic in the dual variable $\boldsymbol{\lambda}$) is negative. So, again, it makes sense to minimize over the primal variable \mathbf{x} and to maximize over the dual variable $\boldsymbol{\lambda}$.

In general, the dual function is $g(\boldsymbol{\lambda}) = \min_{\mathbf{x}} J_A(\mathbf{x}, \boldsymbol{\lambda})$.

However, this may not also be feasible to solve \mathbf{x} explicitly from $\frac{\partial J(\mathbf{x}, \boldsymbol{\lambda})}{\partial \mathbf{x}}$.

Lagrange multipliers do have some similarities with penalty based methods like LASSO and Ridge.

For example, the well-known LASSO objective function $\min_{\boldsymbol{\beta}} \frac{1}{n} \|\mathbf{y} - \mathbf{X} \boldsymbol{\beta}\|_2^2 + \lambda \|\boldsymbol{\beta}\|_1$ is a Lagrangian dual of

$$\min_{\boldsymbol{\beta}} \left\{ \frac{1}{n} \|\mathbf{y} - \mathbf{X} \boldsymbol{\beta}\|_2^2 \right\} \text{ subject to } \|\boldsymbol{\beta}\|_1 \leq t.$$

Similar representations also hold for Ridge or Elastic net.

1 In case of equality constraint:

Consider the problem

$$\min_{\mathbf{x}} f(\mathbf{x}) \text{ such that } \mathbf{Ax} = \mathbf{b} \quad (1)$$

The entire subsection is to solve this problem.

Dual gradient ascent:

The Lagrangian is,

$$J(\mathbf{x}, \boldsymbol{\lambda}) = f(\mathbf{x}) + \boldsymbol{\lambda}^T (\mathbf{Ax} - \mathbf{b})$$

Our objective is to $\max_{\boldsymbol{\lambda}} \min_{\mathbf{x}} J(\mathbf{x}, \boldsymbol{\lambda})$.

The dual gradient ascent repeats for $k = 1, 2, 3, \dots$

$$\begin{aligned} \mathbf{x}^{(k+1)} &= \arg \min_{\mathbf{x}} f(\mathbf{x}) + (\boldsymbol{\lambda}^{(k)})^T \mathbf{Ax} \text{ (only the part that involves } \mathbf{x}) \\ \boldsymbol{\lambda}^{(k+1)} &= \boldsymbol{\lambda}^{(k)} + t_{k+1} (\mathbf{Ax}^{(k+1)} - \mathbf{b}), \text{ (Gradient-based update for } \boldsymbol{\lambda}) \end{aligned}$$

The step size is chosen just like Gradient descent.

Why ‘ascent’? We need to maximize $J(\mathbf{x}, \boldsymbol{\lambda})$ with respect to $\boldsymbol{\lambda}$ and thus, we are applying the gradient-based move to increase $\boldsymbol{\lambda}$, thereby imposing a stronger penalty. Thus, we are ascending along the loss function with respect to $\boldsymbol{\lambda}$.

In order for dual ascent to work, it requires some additional conditions on the function. To avoid such conditions and bring robustness, the following representation was developed.

2 Penalty methods in optimization

Let an optimization problem be:

$$\min_{\mathbf{x}} f(\mathbf{x}) \text{ such that } g_1(\mathbf{x}) \leq 0, g_2(\mathbf{x}) = 0,$$

The above can be rewritten as

$$\min_{\mathbf{x}} f(\mathbf{x}) \text{ such that } g_1(\mathbf{x}) \leq 0, g_2(\mathbf{x}) \leq 0, -g_2(\mathbf{x}) \leq 0,$$

Hence, we can primarily focus on the following problem

$$\min_{\mathbf{x}} f(\mathbf{x}) \text{ such that } \mathbf{g}(\mathbf{x}) \leq 0,$$




Figure 1: Alternative penalty method

where $\mathbf{g}(\mathbf{x}) = \{g_1(\mathbf{x}), g_2(\mathbf{x}), -g_2(\mathbf{x})\}$.

Let for a general case, $\mathbf{g}(\mathbf{x}) = \{g_1(\mathbf{x}), \dots, g_m(\mathbf{x})\}$ i.e. there are m constraints. Then penalty methods typically define a penalty function $P(\mathbf{x})$ such that: 1) $P(\mathbf{x}) = 0$ if $\mathbf{g}(\mathbf{x}) \leq 0$ and $P(\mathbf{x}) > 0$ if $\mathbf{g}(\mathbf{x}) \not\leq 0$, where $\not\leq$ means for at least one of the functions in $\mathbf{g}(\mathbf{x})$ the constraint is not satisfied. For example $P(\mathbf{x}) = \sum_{i=1}^m \max\{0, g_i(\mathbf{x})\}^2$.

We then solve the following unconstrained problem,

$$\min_{\mathbf{x}} f(\mathbf{x}) + \lambda P(\mathbf{x})$$

You can check this note for further details [2]. Under quadratic penalty with equality constraints, even the following loss is also considered as in Figure 1:

3 Augmented Lagrangians and the Method of Multipliers (MM):

The augmented Lagrangian is,

$$J(\mathbf{x}, \boldsymbol{\lambda}) = f(\mathbf{x}) + \boldsymbol{\lambda}^T (\mathbf{Ax} - \mathbf{b}) + \frac{\rho}{2} \|\mathbf{Ax} - \mathbf{b}\|_2^2,$$

where $\rho > 0$ is called the penalty parameter. The augmented Lagrangian can be viewed as the (unaugmented) Lagrangian associated with the problem

$$\min_{\mathbf{x}} f(\mathbf{x}) + \frac{\rho}{2} \|\mathbf{Ax} - \mathbf{b}\|_2^2 \text{ such that } \mathbf{Ax} = \mathbf{b}$$

Applying dual ascent to the modified problem yields the algorithm:

$$\begin{aligned}\mathbf{x}^{(k+1)} &= \arg \min_{\mathbf{x}} f(\mathbf{x}) + \frac{\rho}{2} \|\mathbf{Ax} - \mathbf{b}\|_2^2 + (\boldsymbol{\lambda}^{(k)})^T \mathbf{Ax} \text{ (only the part that involves } \mathbf{x}) \\ \boldsymbol{\lambda}^{(k+1)} &= \boldsymbol{\lambda}^{(k)} + \rho(\mathbf{Ax}^{(k+1)} - \mathbf{b}), \text{ (Gradient-based update for } \boldsymbol{\lambda}),\end{aligned}$$

which is known as the method of multipliers (MM) for solving our original problem in (1). The penalty parameter ρ now comes as the step-length parameter. In general, we increase ρ periodically to strengthen the constraint. Specifically, we set $\rho = \alpha^k$ for some pre-specified α .

4 Very brief on Alternating direction method of multipliers (ADMM)

ADMM is to solve the following optimization problem,

$$\min_{\mathbf{x}, \mathbf{z}} f(\mathbf{x}) + f(\mathbf{z}) \text{ such that } \mathbf{Ax} + \mathbf{Bz} = \mathbf{c}$$

The augmented Lagrangian is,

$$J(\mathbf{x}, \boldsymbol{\lambda}, \rho) = f(\mathbf{x}) + f(\mathbf{z}) + \boldsymbol{\lambda}^T (\mathbf{Ax} + \mathbf{Bz} - \mathbf{c}) + \frac{\rho}{2} \|\mathbf{Ax} + \mathbf{Bz} - \mathbf{c}\|_2^2,$$

Normal MM algorithm is

$$\begin{aligned}(\mathbf{x}^{(k+1)}, \mathbf{z}^{(k+1)}) &= \arg \min_{\mathbf{x}, \mathbf{z}} f(\mathbf{x}) + f(\mathbf{z}) + \frac{\rho}{2} \|\mathbf{Ax} + \mathbf{Bz} - \mathbf{c}\|_2^2 + (\boldsymbol{\lambda}^{(k)})^T (\mathbf{Ax} + \mathbf{Bz}) \text{ (Part that involves } \mathbf{x}, \mathbf{z}) \\ \boldsymbol{\lambda}^{(k+1)} &= \boldsymbol{\lambda}^{(k)} + \rho(\mathbf{Ax}^{(k+1)} + \mathbf{Bz}^{(k+1)} - \mathbf{c}), \text{ (Gradient-based update for } \boldsymbol{\lambda}),\end{aligned}$$

ADMM algorithm is

$$\begin{aligned}\mathbf{x}^{(k+1)} &= \arg \min_{\mathbf{x}} f(\mathbf{x}) + \frac{\rho}{2} \|\mathbf{Ax} + \mathbf{Bz}^{(k)} - \mathbf{c}\|_2^2 + (\boldsymbol{\lambda}^{(k)})^T (\mathbf{Ax}) \text{ (only the part that involves } \mathbf{x}) \\ \mathbf{z}^{(k+1)} &= \arg \min_{\mathbf{z}} f(\mathbf{z}) + \frac{\rho}{2} \|\mathbf{Ax}^{(k+1)} + \mathbf{Bz} - \mathbf{c}\|_2^2 + (\boldsymbol{\lambda}^{(k)})^T (\mathbf{Bz}) \text{ (only the part that involves } \mathbf{z}) \\ \boldsymbol{\lambda}^{(k+1)} &= \boldsymbol{\lambda}^{(k)} + \rho(\mathbf{Ax}^{(k+1)} + \mathbf{Bz}^{(k+1)} - \mathbf{c}), \text{ (Gradient-based update for } \boldsymbol{\lambda}),\end{aligned}$$

The first two steps will find the exact minimizer having others fixed, and then 3rd step will update the $\boldsymbol{\lambda}$. As you may have guessed, we can generalize the loss function to have even more than 2 terms in the objective function. **In addition, we increase ρ gradually, often geometrically.**

Application in distributed computing: ADMM is a popular method for online and distributed optimization on a large scale. Here is an example for LASSO penalty based linear regression.

Say, we have collected, 10000 datapoints. Thus, we have $(y_1, \mathbf{x}_1), \dots, (y_{10000}, \mathbf{x}_{10000})$. Here \mathbf{x}_i 's are p -dimensional and p is also large, say 1 million. Then the standard LASSO regression would solve $\sum_{i=1}^{10000} (y_i - \mathbf{x}_i^T \boldsymbol{\beta})^2 + \theta \|\boldsymbol{\beta}\|_1$. This may be a computational nightmare in terms of both memory and efficiency. Also, this could be a part of streaming data source. In that case, more data will be collected overtime, and we need to refit this giant model again.

Fortunately, ADMM can help us solve this problem using less memory. However, we might loose some accuracy. We break this data into several subsets, say each with size 1000. Then there will be $10000/1000 = 10$ batches of datasets. In this case, the constraint optimization routine will be,

$$\min_{\boldsymbol{\beta}_1, \dots, \boldsymbol{\beta}_{10}, \boldsymbol{\beta}} \sum_{k=1}^{10} \|\mathbf{y}_k - \mathbf{X}_k^T \boldsymbol{\beta}_k\|_2^2 + \theta \|\boldsymbol{\beta}\|_1 \text{ such that } \boldsymbol{\beta}_k = \boldsymbol{\beta} \text{ for all } k = 1, \dots, 10,$$

where \mathbf{y}_k is a block of data points with index $1000(k-1) + 1, \dots, 1000k$ and \mathbf{X}_k is data matrix of dimension $1000 \times p$ with observations $1000(k-1) + 1, \dots, 1000k$. **Here we have batch-specific separate regression coefficients and then a constraint to bring them the same.**

The augmented Lagrangian is,

$$J(\boldsymbol{\beta}_1, \dots, \boldsymbol{\beta}_{10}, \boldsymbol{\lambda}_1, \dots, \boldsymbol{\lambda}_{10}) = \frac{1}{2} \sum_{k=1}^{10} \|\mathbf{y}_k - \mathbf{X}_k^T \boldsymbol{\beta}_k\|_2^2 + \theta \|\boldsymbol{\beta}\|_1 + \sum_{k=1}^{10} \boldsymbol{\lambda}_k^T (\boldsymbol{\beta}_k - \boldsymbol{\beta}) + \frac{\rho}{2} \sum_{k=1}^{10} \|\boldsymbol{\beta}_k - \boldsymbol{\beta}\|_2^2.$$

ADMM algorithm will then run as,

1. Minimize $\boldsymbol{\beta}_k$'s separately and parallelly. There will be a closed form solution for this as we can reformulate the loss for $\boldsymbol{\beta}_k$ like a least square loss.
2. Minimize $\boldsymbol{\beta}$. This will also have a close form solution based on soft-thresholded estimate from the coordinate descent note.
3. Then gradient ascent update for $\boldsymbol{\lambda}_k$'s.

We may attempt to solve this problem in another HW, will see. But similar ideas are applicable for other large scale distributed optimizations and online computing as well. ADMM is also not the only candidate for distributed computing. But ADMM often helps to transform a problem into a convex problem.

In case of online computing, the objective function is,

$$\min_{\beta_1, \beta_2} f(\beta_1) + f(\beta_2) \text{ such that } \beta_1 = \beta_2,$$

where $f(\beta_1)$ is the loss based on old data and $f(\beta_2)$ is the loss of newly acquired data.

ADMM is generally applied with linear and equality constraint. It is possible to extend ADMM for inequality and linear constraint. For non-linear constraint, I am not sure if there is any well established method yet.

If interested to know more about ADMM, do read [1] (https://stanford.edu/~boyd/papers/pdf/admm_distr_stats.pdf) and the slides at https://web.stanford.edu/class/ee364b/lectures/admm_slides.pdf.

4.1 Some interesting applications of ADMM

1. The LAD regression where we minimize $\|\mathbf{y} - \mathbf{X}\beta\|_1$ can be re-written in ADMM form as

$$\min \|\mathbf{z}\|_1 \text{ such that } \mathbf{X}\beta - \mathbf{z} = \mathbf{y}$$

The augmented Lagrangian is, $\|\mathbf{z}\|_1 + \lambda^T(\mathbf{X}\beta - \mathbf{z} - \mathbf{y}) + \frac{\rho}{2}\|\mathbf{X}\beta - \mathbf{z} - \mathbf{y}\|_2^2$.

2. LASSO regression loss is $\|\mathbf{y} - \mathbf{X}\beta\|_2^2 + \lambda\|\beta\|_1$ whose ADMM form is

$$\min \|\mathbf{y} - \mathbf{X}\beta\|_2^2 + \lambda\|\mathbf{z}\|_1 \text{ such that } \beta = \mathbf{z}$$

The augmented Lagrangian is, $\|\mathbf{y} - \mathbf{X}\beta\|_2^2 + \lambda\|\mathbf{z}\|_1 + \alpha^T(\beta - \mathbf{z}) + \frac{\rho}{2}\|\beta - \mathbf{z}\|_2^2$.

3. In case sparse logistic regression, the augmented Lagrangian for the corresponding ADMM becomes $-\sum_{i=1}^n [y_i \log p_i + (1 - y_i) \log(1 - p_i)] + \lambda\|\mathbf{z}\|_1 + \alpha^T(\beta - \mathbf{z}) + \frac{\rho}{2}\|\beta - \mathbf{z}\|_2^2$, where $p_i = \frac{1}{1 + \exp(\mathbf{x}_i^T \beta)}$

For the sparse logistic regression case, the complete ADMM algorithm steps are:

$$\begin{aligned}\beta^{(k+1)} &= \arg \min_{\beta} - \sum_{i=1}^n [y_i \log p_i + (1 - y_i) \log(1 - p_i)] + \alpha^T (\beta - \mathbf{z}) + \frac{\rho}{2} \|\beta - \mathbf{z}\|_2^2 \\ \mathbf{z}^{(k+1)} &= \arg \min_{\mathbf{z}} \alpha^T (\beta - \mathbf{z}) + \frac{\rho}{2} \|\beta - \mathbf{z}\|_2^2 = S_{\lambda/\rho}(\beta + \alpha/\rho), \text{ Soft-thresholded close-form minimizer} \\ \alpha^{(k+1)} &= \alpha^{(k)} + \rho(\beta^{(k+1)} - \mathbf{z}^{(k+1)}), \text{ (Gradient-based update for } \alpha),\end{aligned}$$

Increase ρ , possibly geometrically, as before.

5 Software packages for optimization

Here are options, `optim` in R or `scikit-optimize` in Sklearn of python (https://scikit-optimize.github.io/dev//_downloads/scikit-optimize-docs.pdf). In Pytorch, `torch.optim` is good.

In Cpp, `ensmallen` library is awesome <https://www.ensmallen.org/docs.html>. It also has a list of optimization functions along with citations. Now `RcppEnsmallen` (<https://cran.r-project.org/web/packages/RcppEnsmallen/vignettes/using-rcppensmallen.html>) is available to use this library R, and it also supports the vectorized efficient functions from `RcppArmadillo`, using `armadillo` Cpp library (<https://arma.sourceforge.net/docs.html>).

5.1 Automatic differentiation

For R, there are two options: 1) <http://www.pqr-project.org/>, and 2) use C++ library `CppAD` within `Rcpp`. (https://mkomod.github.io/user_tn/technical_note.html) There is another (super-slow) option using R package `pracma`.

For python, there are several options: 1) `Pytorch` and `tensorflow` have their own (<https://pytorch.org/docs/stable/autograd.html> is good), but you need the data in `tensor` formatted data structure, 2) `NumPy` is one of the most popular package which is like `base` package R and there is a new package called `jax` (also requires `jaxlib`) that uses `NumPy` data structure and have functions for automatic differentiation.

References

- [1] Stephen Boyd, Neal Parikh, Eric Chu, Borja Peleato, Jonathan Eckstein, et al. Distributed optimization and statistical learning via the alternating direction method of multipliers. *Foundations and Trends® in Machine learning*, 3(1):1–122, 2011.

- [2] Robert M Freund. Penalty and barrier methods for constrained optimization. *Lecture Notes, Massachusetts Institute of Technology*, 2004.